



UPPSALA
UNIVERSITET

Parallell and Distributed Programming **Assignment 1**

Authors:

Aikaterini Manousidou

Moa Li Lekander

Uppsala

April 20, 2021

Contents

1	Introduction	1
1.1	Theory	1
2	Parallelization	2
3	Verification of Correctness	3
4	Performance Experiments	4
5	Discussion	7
6	Appendix	9
6.1	Serial Code	9
6.2	Parallelized code	11
6.3	Performance Verification and Evaluation Code	14

1 Introduction

A commonly used tool within image processing and computer simulations, is the operation of stencil applications. In this assignment a serial code is provided that initiates a one dimensional array of function values $f(x_i) = \sin(\frac{2\pi i}{N})$, where N is the total number of function values in the array and i is the index of the value. Then, a stencil is applied on every element in the array. The stencil in this assignment is defined so that the two adjacent values on either side of the element in question are added together to obtain the new value, according to equation (1). In (1), $f(x_i)$ represents the element the stencil is applied to and $f(x_{i-j})$ are its adjacent elements to the left and $f(x_{i+j})$ the adjacent elements to the right. This stencil is used to approximate the first derivative of the function values in the array.

$$\frac{1}{12h} \cdot f(x_{i-2}) - \frac{8}{12h} \cdot f(x_{i-1}) + 0 \cdot f(x_i) + \frac{8}{12h} \cdot f(x_{i+1}) - \frac{1}{12h} \cdot f(x_{i+2}) \quad (1)$$

The aim of this assignment is to parallelize the given serial code using the *Message Passing Interface* (MPI). The number of function values and the number of times the stencil is to be applied are provided as input to the program. The boundary values of the array are periodic. Additionally, the implementation of the parallelized code assumes that the number of function values are divisible by the specified number of processes and that the number of values each process works on is larger than the width of the stencil. After the parallelized implementation is made its functionality will be verified and its performance evaluated by examining both the strong and weak scaling of the program.

1.1 Theory

A parallelized program can be evaluated with two different types of scalings, a strong and a weak scaling. A strong scaling analysis implies that the size of the program is kept constant while increasing the number of processes. A tool for the evaluation of the performance of a parallelized program is calculating the strong scaling speed-up, as seen in equation (2), where N is the size of the problem, p is the number of processes, S is the speed-up and T_{serial} and $T_{parallel}$ are the timings for the serial and parallelized codes, respectively. Ideally, the speed-up $S(N, p)$ should be equal to p . In other words, it should be linear and equal to the number of processes used.[1]

$$S(N, p) = \frac{T_{serial}(N)}{T_{parallel}(N, p)} \quad (2)$$

On the other hand, a weak scaling analysis indicates that the problem size increases relative to the number of processes, so that the workload for each process is kept constant. The speed-up for this type of scaling can be seen in equation (3).[2][3]

$$S(N, p) = \frac{T_{serial}(N)}{T_{parallel}(N, p)} \times p \quad (3)$$

Another tool that is widely used to evaluate parallel programs is the parallel efficiency, which can also be perceived as the speed-up per process and can be found in equation (4).

$$E(N, p) = \frac{S(N, p)}{p} = \frac{T_{serial}(N)}{p \times T_{parallel}(N, p)} \quad (4)$$

As a result, a program that retains constant efficiency while the problem size remains the same is called strongly scalable, although a program that retains constant efficiency while the problem size increases is called weakly scalable.[1]

2 Parallelization

The computations were parallelized by equally dividing the array values among the processes. Firstly, the process with rank zero fills the input array with function values, thereafter, the array is divided among the processes by the `MPI_Scatter` function. This function was used because it provides an efficient way of splitting the array into parts and sending each of the them to different processes. In every process, the given part of the array is stored in a local input array which has a total of two extra elements (initially empty) on each end. This was implemented for the purpose of handling the edges of the local arrays correctly.

A for-loop is used to repeatedly apply the stencil to the arrays. In the beginning of every loop iteration, the four extra elements of the local input arrays are filled with values by using the non-blocking send and receive functions `MPI_Isend` and `MPI_Irecv` together with `MPI_Wait`. These are used so that the function values needed to apply the stencil to the edge elements of the local array are obtained. Hence, every process communicates with the processes that store the adjacent parts of the array as it sends and receives the edge values. When all necessary function values are provided, the stencil application occurs and the new values are stored in a local output vector. Thereafter, a call to the function `MPI_Barrier` is made. This barrier function is used for the purpose of synchronizing the processes, in order to make sure that all processes have finished computing their new values for the given step before they all move on to the next iteration and apply the stencil again. Therefore, when all processes have called the barrier function the program will continue and the local input vector will be filled with the newly computed values stored in the local output vector before moving on to the next loop iteration.

When the stencil has been applied the given number of steps, a call to the function `MPI_Gather` is made by the process with rank zero, in order to store all the final values in an output array. The maximum stencil application time for the processes is computed using the function `MPI_Reduce`, and finally all allocated memory is being freed before the program terminates. Both the serial and the parallelized programs can be found in the Appendix sections 6.1 and 6.2, respectively.

3 Verification of Correctness

In order to make sure that the functionality of the serial program was kept as the parallelized version was implemented, a verification of correctness was made. This verification was made by comparing the values in the produced output arrays obtained from the serial and parallelized programs, for different input values. Different number of processes were also used in the tests in order to verify that the program was working for different number of processes. Figure 1 presents the result of these measurements, where it can be observed that the computed output values are the same for both programs. In the two top plots in Figure 1 the output strongly resembles a negative sine curve, which is the expected result since the 10th derivative of a sine curve is a negative sine curve. In the same way, the output in the bottom plot in Figure 1 resembles a cosine curve, which corresponds to the 5th derivative of a sine curve. From these results, it could be concluded that the functionality of the serial program was preserved in the parallelized version.

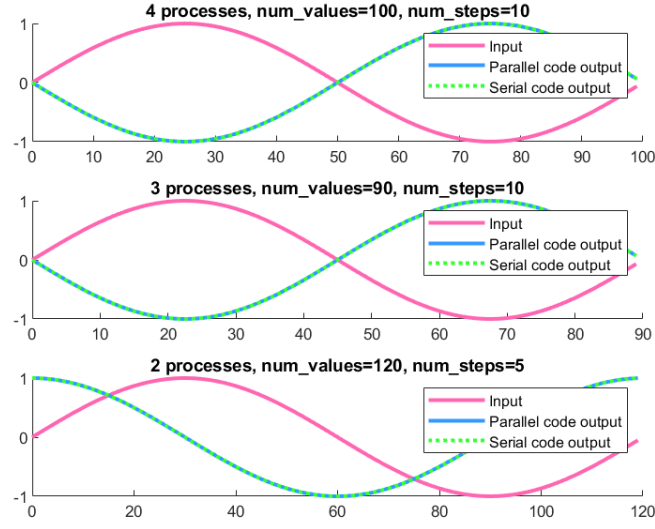


Figure 1: Input function values together with output function values for both the serial and the parallelized programs, for different input parameters and number of processes used for the parallelized code

4 Performance Experiments

In order to evaluate the performance of the parallelized code, a series of experiments was carried out. These experiments were performed on one of the scientific Linux hosts, **gullviva**. This specific computer acquires an AMD Opteron(TM) Processor 6274 with 2 sockets and 8 cores per socket, resulting in a total of 16 cores. This implies that the best and highest performance should be achieved when using 16 processes. Thereafter, the compiler Open MPI with version 1.8.1 was used to compile both the serial and parallelized programs.

First, a strong scaling analysis was performed. This implicated measuring the stencil application time for a problem of a fixed size but for an increasing number of processes. The number of function values that was chosen for this experiment was 1.6×10^7 in order to maximize the amount of processes this number was divisible with. The number of steps was set to 100 and the amounts of processes tested for the parallelized code were 2, 4, 5, 8, 10, and 16. After obtaining the stencil application times, the speed-up was calculated with the help of equation (2) and Figure 2 was generated, illustrating the relations between the number of processes, stencil application time and speed-up. Figure 2 also portrays the ideal values of the stencil applications timings and the ideal speed-up. The numeric values of the measured timings can be found in Table 1.

Table 1: Table over the number of processes, the stencil application timings acquired during the strong scaling analysis and the calculated speed-up. All values have been rounded to five significant digits.

Processes	Stencil application time (s)	Speed-up
1	76.0257	1.0000
2	73.3039	1.0371
4	32.0735	2.3704
5	37.9536	2.0031
8	23.4678	3.2396
10	25.7712	2.9500
16	19.0992	3.9806

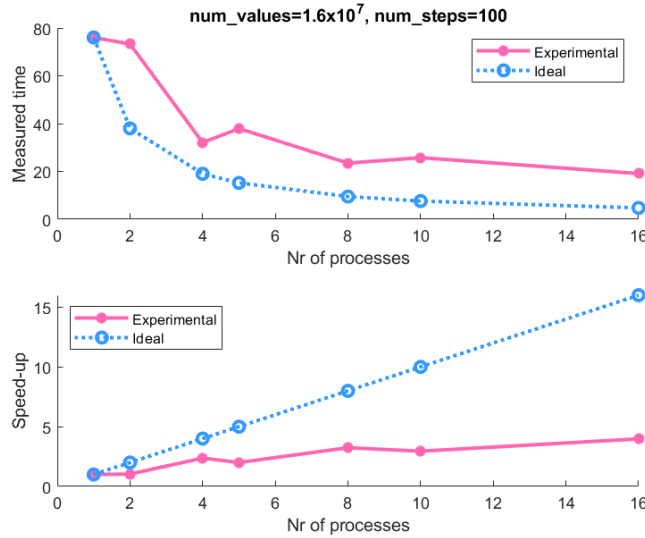


Figure 2: Plot over the strong scaling analysis, illustrating the relations between the number of processes and the measured timings and the number of processes and the speed-up, respectively.

Then, a weak scaling analysis was performed. This implied measuring the stencil application time for a number of values of $1000000 \times p$, where p depicts the number of processes. The number of steps was, once again, set to 100 and the numbers of processes tested for the parallelized code were increased from 2 to 16, with increments of 1. After measuring the stencil application timings, the speed-up was determined with the help of equation (3). Figure 3 was thereupon produced, depicting the relations between the number of processes, stencil application time and speed-up. In Figure 3 the ideal values of the stencil applications timings and the the speed-up are also represented. The numeric values of the measured timings can be found in Table 2.

Table 2: Table over the number of processes, the stencil application timings acquired during the weak scaling analysis and the calculated speed-up. All values have been rounded to five significant digits.

Processes	Problem size	Stencil application time (s)	Speed-up
1	1×10^6	0.4465	1.0000
2	2×10^6	0.7506	1.1898
3	3×10^6	0.7562	1.7716
4	4×10^6	0.7712	2.3160
5	5×10^6	1.0157	2.1983
6	6×10^6	1.0515	2.5480
7	7×10^6	1.0625	2.9419
8	8×10^6	1.0943	3.2645
9	9×10^6	1.3855	2.9007
10	10×10^6	1.4634	3.0515
11	11×10^6	1.4390	3.4135
12	12×10^6	1.4444	3.7099
13	13×10^6	1.7913	3.2407
14	14×10^6	1.8479	3.3830
15	15×10^6	1.8560	3.6089
16	16×10^6	2.3554	3.0333

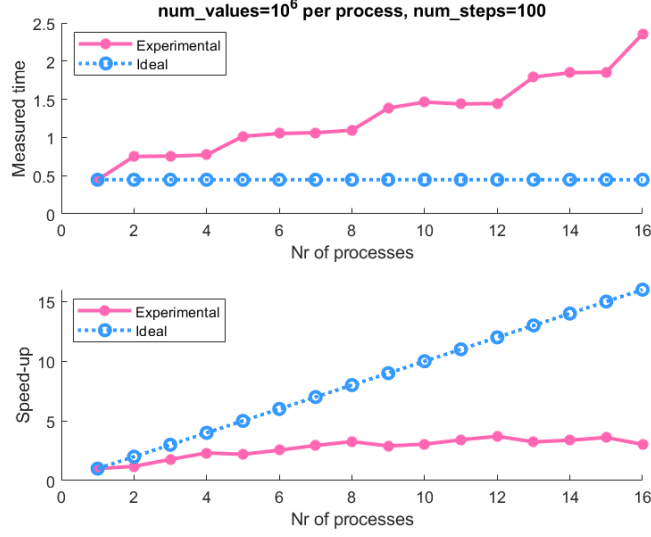


Figure 3: Plot over the weak scaling analysis, illustrating the relations between the number of processes and the measured time and the number of processes and the speed-up, respectively.

Finally, Figure 4 was produced by calculating the efficiency with the help of equation (4). The code for the verification and evaluation of the parallelized program can be found in the Appendix section 6.3.

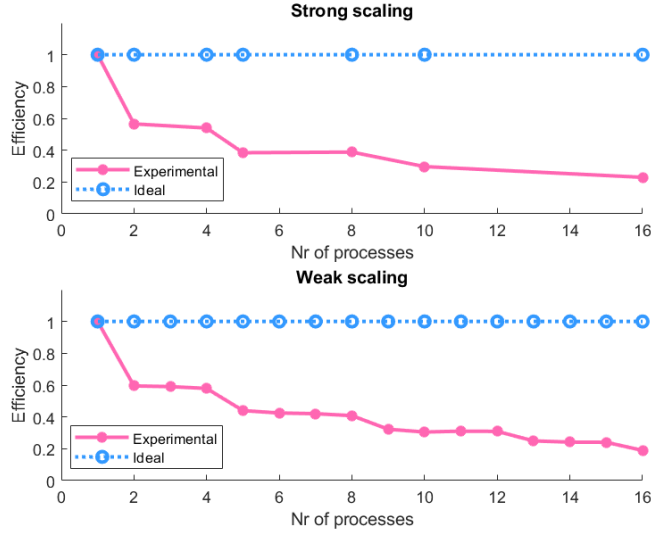


Figure 4: Plot over the efficiencies of the strong and weak scaling analysis.

5 Discussion

From the results presented in Figure 2, it can be concluded that the measured timings seem to follow the trend line of the ideal case. However, there is a difference in the time for the measured values as they are slightly larger than the ideal values, which thereby results in a smaller speedup. This implies that the parallelized program does not appear to be strongly scalable.

When analyzing the weak scaling presented in Figure 3 it can be noted that the measured time increases as the amount of processes involved increase. The curve seems to have a somewhat linear appearance, and as a result the speedup seems to be constant throughout the different number of processes. By comparing to the ideal case, the measured curves seem to have different slopes. This implicates that the parallelized code does not seem to be weakly scalable.

From Figure 4 it can be concluded that the final and parallelized program does not appear to follow the ideal cases, on the contrary it seems to decrease in both types of scalings. This, once again, confirms the conclusions drawn from Figures 2 and 3 that the program is neither strongly scalable nor weakly scalable. On the other hand, it can be noted that the program's best performance occurs when the maximum amount of processes is being used, which is as expected.

References

- [1] P. Pacheco, *An Introduction to Parallel Programming*. International series of monographs on physics, Morgan Kaufmann, 2011.
- [2] F. Wermelinger, Computational Science and Engineering Laboratory, “[Strong and Weak Scaling](#).”
- [3] Research Computing, Office of Information Technology, University of Colorado Boulder, “[Scaling Analysis](#).”

6 Appendix

6.1 Serial Code

```
1  #define PI 3.14159265358979323846
2  #include <mpi.h>
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <math.h>
6
7  int main(int argc, char **argv) {
8      if (3 != argc) {
9          printf("Usage: stencil num_values num_steps\n");
10         return 0;
11     }
12     int num_values = atoi(argv[1]);
13     int num_steps = atoi(argv[2]);
14     int rank, size;
15
16     MPI_Init(&argc, &argv);
17     MPI_Comm_size(MPI_COMM_WORLD, &size); /* Get the number of processors */
18     MPI_Comm_rank(MPI_COMM_WORLD, &rank); /* Get my number */
19
20     // Generate values for stencil operation
21     double *input=(double*)malloc(num_values*sizeof(double));
22     double h = 2.0*PI/num_values;
23     for (int i=0; i<num_values; i++) input[i]=sin(h*i);
24
25     // Stencil values
26     const int STENCIL_WIDTH = 5;
27     const int EXTENT = STENCIL_WIDTH/2;
28     const double STENCIL[] = {1.0/(12*h), -8.0/(12*h), 0.0, 8.0/(12*h), -1.0/(12*h)};
29
30     // Start timer
31     double start = MPI_Wtime();
32
33     // Allocate data for result
34     double *output =(double*)malloc(num_values*sizeof(double));
35
36     // Print input values
37     /*if(rank == 0){
38         printf("input:\n");
39         for(int i=0; i<num_values; i++){
40             printf("%f ", input[i]);
41         }
42         printf("\n");
43     }*/
44
45     // Repeatedly apply stencil
46     for (int s=0; s<num_steps; s++) {
47         // Apply stencil on left boundary with periodic cond
48         for (int i=0; i<EXTENT; i++) {
49             double result = 0;
50             for (int j=0; j<STENCIL_WIDTH; j++) {
51                 int index = (i - EXTENT + j + num_values) % num_values;
52                 result += STENCIL[j] * input[index];
53             }
54             output[i] = result;
55         }
56
57         // Apply stencil on inner points
58         for (int i=EXTENT; i<num_values-EXTENT; i++) {
59             double result = 0;
```

```

60         for (int j=0; j<STENCIL_WIDTH; j++) {
61             int index = i - EXTENT + j;
62             result += STENCIL[j] * input[index];
63         }
64         output[i] = result;
65     }
66
67     // Apply stencil on right boundary with periodic cond
68     for (int i=num_values-EXTENT; i<num_values; i++) {
69         double result = 0;
70         for (int j=0; j<STENCIL_WIDTH; j++) {
71             int index = (i - EXTENT + j) % num_values;
72             result += STENCIL[j] * input[index];
73         }
74         output[i] = result;
75     }
76
77     // Swap input and output
78     if (s < num_steps-1) {
79         double *tmp = input;
80         input = output;
81         output = tmp;
82     }
83 }
84
85 // Stop timer
86 double my_execution_time = MPI_Wtime() - start;
87 printf("%f\n", my_execution_time);
88
89 //Print output vector
90 /*if(rank == 0){
91     printf("\noutput:\n");
92     for(int i=0; i<num_values; i++){
93         printf("%f ", output[i]);
94     }
95     printf("\n");
96 }*/
97
98 // Write output to file
99 /*FILE *file=fopen("output_old.txt","w");
100 for (int i = 0; i < num_values; i++)
101     fprintf(file, "%f \n", output[i]);
102 fclose(file);
103 */
104
105 // Clean up
106 free(input);
107 free(output);
108 MPI_Finalize();
109 return 0;
110 }

```

6.2 Parallelized code

```
1  #define PI 3.14159265358979323846
2  #include <mpi.h>
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <math.h>
6
7  int main(int argc, char **argv) {
8      if (3 != argc) {
9          printf("Usage: stencil num_values num_steps\n");
10         return 0;
11     }
12     int num_values = atoi(argv[1]);
13     int num_steps = atoi(argv[2]);
14     int rank, size;
15     double execution_time, total_start_time, start_time, max_time, total_time;
16
17     MPI_Status status;
18     MPI_Request request;
19
20     MPI_Init(&argc, &argv);
21     MPI_Comm_size(MPI_COMM_WORLD, &size); /* Get the number of processors */
22     MPI_Comm_rank(MPI_COMM_WORLD, &rank); /* Get my number */
23     int chunk = num_values/size; // number of values for every process
24
25     // Stencil values
26     const int STENCIL_WIDTH = 5;
27     const int EXTENT = STENCIL_WIDTH/2;
28     double h = 2.0*PI/num_values;
29     const double STENCIL[] = {1.0/(12*h), -8.0/(12*h), 0.0, 8.0/(12*h), -1.0/(12*h)};
30
31     // Check assumptions
32     if(num_values%size != 0) {
33         if(rank==0) printf("ERROR: num_values must be divisible by number of processes!\n");
34         exit(0);
35     } else if(chunk < STENCIL_WIDTH){
36         if(rank==0) printf("ERROR: Too many processors for given num_values!\n");
37         exit(1);
38     }
39
40     // Allocate data for input and output
41     double *input=(double*)malloc(num_values*sizeof(double));
42     double *output =(double*)malloc(num_values*sizeof(double));
43     double *local_input =(double *)malloc((chunk+EXTENT*2)*sizeof(double));
44     double *local_output =(double *)malloc(chunk*sizeof(double));
45
46     if(rank==0){
47         // Generate values for stencil operation
48         for (int i=0; i<num_values; i++) input[i]=sin(h*i);
49
50         // Print input values
51         /*printf("input:\n");
52         for(int i=0; i<num_values; i++){
53             printf("%f ", input[i]);
54         }
55         printf("\n");*/
56
57         // Write input values to file
58         /*
59         FILE *file1=fopen("input.txt","w");
60         for (int i = 0; i < num_values; i++)
61             fprintf(file1, "%f \n", input[i]);
```

```

62     fclose(file1);*/
63
64     total_start_time = MPI_Wtime();
65 }
66
67 MPI_Scatter(&input[0], chunk, MPI_DOUBLE, &local_input[EXTENT], chunk, MPI_DOUBLE, 0, ...
        MPI_COMM_WORLD);
68
69 // Start timer
70 start_time = MPI_Wtime();
71
72 // Repeatedly apply stencil
73 for (int s=0; s<num_steps; s++) {
74     if(rank==0){
75         MPI_Isend(&local_input[chunk+EXTENT-2], 2, MPI_DOUBLE, 1, 300+1, ...
            MPI_COMM_WORLD, &request);
76         MPI_Isend(&local_input[2], 2, MPI_DOUBLE, size-1, 100+(size-1), ...
            MPI_COMM_WORLD, &request);
77
78         MPI_Irecv(&local_input[chunk+EXTENT], 2, MPI_DOUBLE, 1, 100+rank, ...
            MPI_COMM_WORLD, &request);
79         MPI_Wait(&request, &status);
80
81         MPI_Irecv(&local_input[0], 2, MPI_DOUBLE, size-1, 300+0, MPI_COMM_WORLD, ...
            &request);
82         MPI_Wait(&request, &status);
83     } else if(rank==size-1){
84         MPI_Isend(&local_input[chunk+EXTENT-2], 2, MPI_DOUBLE, 0, 300+0, ...
            MPI_COMM_WORLD, &request);
85         MPI_Isend(&local_input[2], 2, MPI_DOUBLE, rank-1, 100+(rank-1), ...
            MPI_COMM_WORLD, &request);
86
87         MPI_Irecv(&local_input[chunk+EXTENT], 2, MPI_DOUBLE, 0, 100+rank, ...
            MPI_COMM_WORLD, &request);
88         MPI_Wait(&request, &status);
89         MPI_Irecv(&local_input[0], 2, MPI_DOUBLE, rank-1, 300+rank, MPI_COMM_WORLD, ...
            &request);
90         MPI_Wait(&request, &status);
91     } else {
92         MPI_Isend(&local_input[2], 2, MPI_DOUBLE, rank-1, 100+(rank-1), ...
            MPI_COMM_WORLD, &request);
93         MPI_Isend(&local_input[chunk+EXTENT-2], 2, MPI_DOUBLE, rank+1, 300+(rank+1), ...
            MPI_COMM_WORLD, &request);
94
95         MPI_Irecv(&local_input[0], 2, MPI_DOUBLE, rank-1, 300+rank, MPI_COMM_WORLD, ...
            &request);
96         MPI_Wait(&request, &status);
97         MPI_Irecv(&local_input[chunk+EXTENT], 2, MPI_DOUBLE, rank+1, 100+rank, ...
            MPI_COMM_WORLD, &request);
98         MPI_Wait(&request, &status);
99     }
100 }
101
102 // Apply stencil
103 for (int i=EXTENT; i<chunk+EXTENT; i++) {
104     double result = 0;
105     for (int j=0; j<STENCIL_WIDTH; j++) {
106         int index = i - EXTENT + j;
107         result += STENCIL[j] * local_input[index];
108     }
109     local_output[i-EXTENT] = result;
110 }
111
112 MPI_Barrier(MPI_COMM_WORLD); // synchronize processes
113
114 // Swap input and output

```

```

114         if (s < num_steps-1) {
115             for(int i=0; i<chunk; i++){
116                 local_input[i+EXTENT] = local_output[i];
117             }
118         } else {
119             double execution_time = MPI_Wtime()-start_time; // stop timer
120             MPI_Gather(&local_output[0], chunk, MPI_DOUBLE, &output[0], chunk, MPI_DOUBLE, ...
121                     0, MPI_COMM_WORLD);
122             if (rank == 0) total_time = MPI_Wtime()-total_start_time; // stop timer
123         }
124     }
125     // Print output values
126     /*if(rank==0){
127         printf("output:\n");
128         for(int i=0; i<num_values; i++){
129             printf("%f ", output[i]);
130         }
131         printf("\n");
132     }*/
133
134     execution_time = MPI_Wtime()-start_time; // stop timer
135
136     // Find max time among processors
137     MPI_Reduce(&execution_time, &max_time, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
138
139     // Display timing results
140     if(rank==0){
141         printf("Maximum stencil application time: %f\n", max_time);
142         printf("Total time: %f\n", total_time);
143     }
144
145     /*
146     // Write to file
147     FILE *file=fopen("output.txt","w");
148     for (int i = 0; i < num_values; i++)
149         fprintf(file, "%f \n", output[i]);
150     fclose(file);
151     */
152
153     // Clean up
154     free(input);
155     free(output);
156     free(local_input);
157     free(local_output);
158     MPI_Finalize();
159     return 0;
160 }

```

6.3 Performance Verification and Evaluation Code

```
1  close all;
2  clear all;
3
4  %----- PDP ASSIGNMENT 1 -----
5
6  %----- Evaluation of output -----
7
8  %Importing data from text-file and putting it in an array
9  input100= importdata ('input100.txt');
10 output100= importdata ('output2_100.txt');
11 output_old100= importdata ('output_old100.txt');
12 time100 = 0:1:99;
13
14 input90= importdata ('input90.txt');
15 output90= importdata ('output2_90.txt');
16 output_old90= importdata ('output_old90.txt');
17 time90 = 0:1:89;
18
19 input120= importdata ('input120.txt');
20 output120= importdata ('output2_120.txt');
21 output_old120= importdata ('output_old120.txt');
22 time120 = 0:1:119;
23
24 figure(1)
25 subplot(3,1,1);
26 hold on
27 plot(time100,input100, 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
28 plot(time100,output100, 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
29 plot(time100,output_old100, ':', 'Color', [51/255 255/255 51/255], 'LineWidth', 2);
30 legend('Input', 'Parallel code output', 'Serial code output');
31 title('4 processes, num\_values=100, num\_steps=10');
32 hold off
33 subplot(3,1,2);
34 hold on
35 plot(time90,input90, 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
36 plot(time90,output90, 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
37 plot(time90,output_old90, ':', 'Color', [51/255 255/255 51/255], 'LineWidth', 2);
38 legend('Input', 'Parallel code output', 'Serial code output');
39 title('3 processes, num\_values=90, num\_steps=10');
40 hold off
41 subplot(3,1,3);
42 hold on
43 plot(time120,input120, 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
44 plot(time120,output120, 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
45 plot(time120,output_old120, ':', 'Color', [51/255 255/255 51/255], 'LineWidth', 2);
46 legend('Input', 'Parallel code output', 'Serial code output');
47 title('2 processes, num\_values=120, num\_steps=5');
48 hold off
49
50 %----- Strong Scaling -----
51
52 cores1 = [1, 2, 4, 5, 8, 10, 16];
53
54 % Measured times
55
56 % num_values = 16 000 000
57 % num_steps = 100
58
59 % T1 = 6.795839; % 1 core
60 % time1 = zeros(1,length(cores1));
61 % time1(1) = T1;
```



```

62 % time1(2) = 6.01033; % 2 cores
63 % time1(3) = 3.149910; % 4 cores
64 % time1(4) = 3.538132; % 5 cores
65 % time1(5) = 2.188927; % 8 cores
66 % time1(6) = 2.288942; % 10 cores
67 % time1(7) = 1.850582; % 16 cores
68
69 % num_values = 160 000 000
70 % num_steps = 100
71
72 T1 = 76.025657; % 1 core
73 time1 = zeros(1,length(cores1));
74 time1(1) = T1;
75 time1(2) = 73.303855; % 2 cores
76 time1(3) = 32.073502; % 4 cores
77 time1(4) = 37.953636; % 5 cores
78 time1(5) = 23.467748; % 8 cores
79 time1(6) = 25.771217; % 10 cores
80 time1(7) = 19.099210; % 16 cores
81
82 % Calculating speed-up, efficiency and ideal time
83 speedup1 = zeros(1,length(cores1));
84 efficiency1 = zeros(1,length(cores1));
85 t1 = zeros(1,length(cores1));
86 for i=1:length(cores1)
87     speedup1(i) = T1/time1(i);
88     efficiency1(i) = speedup1(i)/cores1(i);
89     t1(i) = T1/cores1(i);
90 end
91
92 % Generating figures
93 figure(2)
94 subplot(2,1,1)
95 hold on
96 plot(cores1, time1, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
97 plot(cores1, t1, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
98 legend('Experimental', 'Ideal');
99 title('num\_values=1.6x10^7, num\_steps=100');
100 xlabel('Nr of processes');
101 ylabel('Measured time');
102 hold off
103 subplot(2,1,2)
104 hold on
105 plot(cores1,speedup1, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
106 plot(cores1, cores1, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
107 legend('Experimental', 'Ideal');
108 xlabel('Nr of processes');
109 ylabel('Speed-up');
110 hold off
111
112 %----- Weak Scaling -----
113
114 cores2 = 1:1:16;
115
116 % Measured times
117
118 % 1000 num_values / process
119 % num_steps = 100
120
121 % T2 = 0.000323; % 1 core
122 % time2 = zeros(1,length(cores2));
123 % time2(1) = T2;
124 % time2(2) = 0.001286; % 2 cores
125 % time2(3) = 0.001558; % 3 cores
126 % time2(4) = 0.001879; % 4 cores

```

```

127 % time2(5) = 0.001912; % 5 cores
128 % time2(6) = 0.002038; % 6 cores
129 % time2(7) = 0.002130; % 7 cores
130 % time2(8) = 0.002668; % 8 cores
131 % time2(9) = 0.002908; % 9 cores
132 % time2(10) = 0.002942; % 10 cores
133 % time2(11) = 0.002956; % 11 cores
134 % time2(12) = 0.003156; % 12 cores
135 % time2(13) = 0.002898; % 13 cores
136 % time2(14) = 0.003546; % 14 cores
137 % time2(15) = 0.003381; % 15 cores
138 % time2(16) = 0.003897; % 16 cores
139
140 % 1 000 000 num_values / process
141 % num_steps = 100
142
143 T2 = 0.446543;
144 time2 = zeros(1,length(cores2));
145 time2(1) = T2;
146 time2(2) = 0.750622; % 2 cores
147 time2(3) = 0.756170; % 3 cores
148 time2(4) = 0.771220; % 4 cores
149 time2(5) = 1.015651; % 5 cores
150 time2(6) = 1.051516; % 6 cores
151 time2(7) = 1.062520; % 7 cores
152 time2(8) = 1.094302; % 8 cores
153 time2(9) = 1.385474; % 9 cores
154 time2(10) = 1.463371; % 10 cores
155 time2(11) = 1.438969; % 11 cores
156 time2(12) = 1.444379; % 12 cores
157 time2(13) = 1.791291; % 13 cores
158 time2(14) = 1.847930; % 14 cores
159 time2(15) = 1.856020; % 15 cores
160 time2(16) = 2.355385; % 16 cores
161
162 % Calculating speed-up, efficiency and ideal time
163 speedup2 = zeros(1,length(cores2));
164 efficiency2 = zeros(1,length(cores2));
165 t2 = T2*ones(1,length(cores2));
166 for i=1:length(cores2)
167     speedup2(i) = T2/time2(i)*cores2(i);
168     efficiency2(i) = speedup2(i)/cores2(i);
169 end
170
171 % Generating figures
172 figure(3)
173 subplot(2,1,1)
174 hold on
175 plot(cores2, time2, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
176 plot(cores2, t2, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
177 legend('Experimental', 'Ideal');
178 title('num\_values=10^6 per process, num\_steps=100');
179 xlabel('Nr of processes');
180 ylabel('Measured time');
181 hold off
182 subplot(2,1,2)
183 hold on
184 plot(cores2,speedup2, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
185 plot(cores2, cores2, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
186 legend('Experimental', 'Ideal');
187 xlabel('Nr of processes');
188 ylabel('Speed-up');
189 hold off
190
191 % Calculating ideal efficiency

```

```

192 e1 = ones(1,length(cores1));
193 e2 = ones(1,length(cores2));
194
195 figure(4)
196 subplot(2,1,1)
197 hold on
198 plot(cores1, efficiency1, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
199 plot(cores1, e1, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
200 ylim([0 1.2])
201 legend('Experimental', 'Ideal');
202 title('Strong scaling');
203 xlabel('Nr of processes');
204 ylabel('Efficiency');
205 hold off
206 subplot(2,1,2)
207 hold on
208 plot(cores2, efficiency2, '-*', 'Color', [255/255 102/255 178/255], 'LineWidth', 2);
209 plot(cores2, e2, ':o', 'Color', [51/255 153/255 255/255], 'LineWidth', 2);
210 ylim([0 1.2])
211 legend('Experimental', 'Ideal');
212 title('Weak scaling');
213 xlabel('Nr of processes');
214 ylabel('Efficiency');
215 hold off

```